



# Contents

- **Transaction Concept**
- **Transaction State**
- **Implementation of Atomicity  
and Durability**
- **Concurrent Executions**
- **Serializability**
- **Recoverability**



# Serializability

- Basic Assumption – Each transaction preserves database consistency.
- Thus serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is **serializable** if it is equivalent to a **serial schedule**. Different forms of schedule equivalence give rise to the notions of:
  1. **conflict serializability**
  2. **view serializability**
- We ignore operations other than **read** and **write** instructions, and we assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.

Our simplified schedules consist of only **read** and **write** instructions



# Conflict Serializability

- Instructions  $l_i$  and  $l_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $l_i$  and  $l_j$ , and at least one of these instructions wrote  $Q$ .
  1.  $l_i = \mathbf{read}(Q)$ ,  $l_j = \mathbf{read}(Q)$ .  $l_i$  and  $l_j$  don't conflict.
  2.  $l_i = \mathbf{read}(Q)$ ,  $l_j = \mathbf{write}(Q)$ . They conflict.
  3.  $l_i = \mathbf{write}(Q)$ ,  $l_j = \mathbf{read}(Q)$ . They conflict
  4.  $l_i = \mathbf{write}(Q)$ ,  $l_j = \mathbf{write}(Q)$ . They conflict
- Intuitively, a conflict between  $l_i$  and  $l_j$  forces a (logical) temporal order between them. If  $l_i$  and  $l_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.



# Conflict Serializability (Cont.)

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**.
- We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule
- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
<b>read(<math>Q</math>)</b>	
	<b>write(<math>Q</math>)</b>

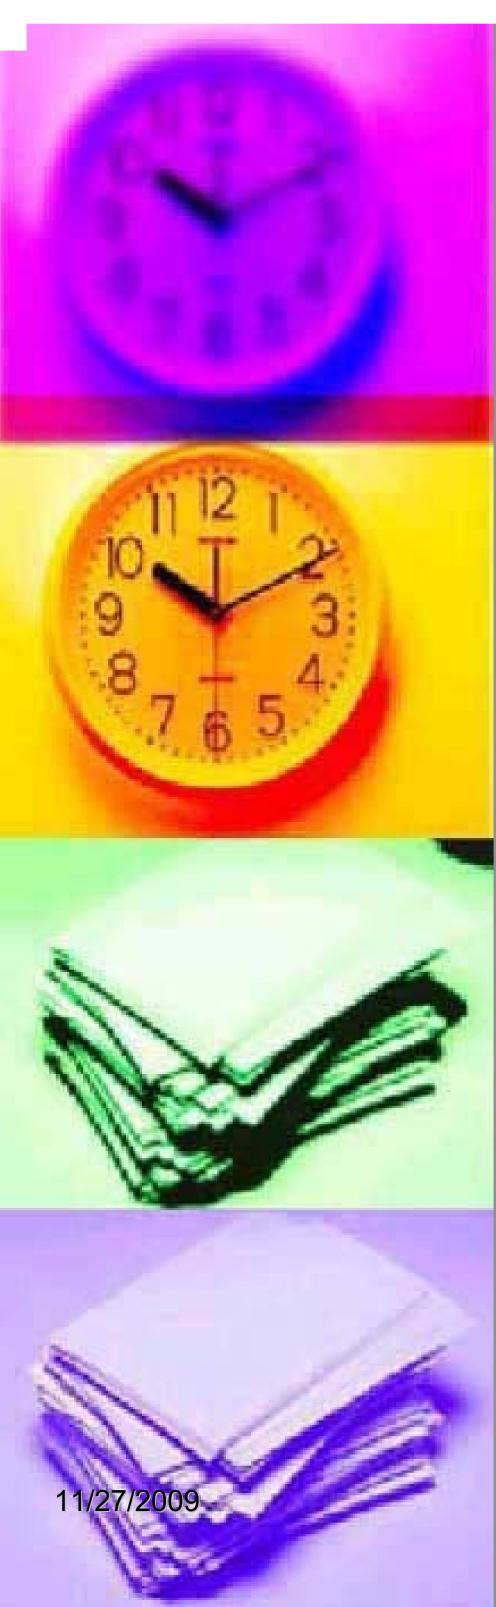
We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .



# Conflict Serializability (Cont.)

- Schedule1 below can be transformed into Schedule5 , a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions.  
Therefore Schedule1 is conflict serializable.

$T_1$	$T_2$
read( $A$ ) write( $A$ )	
read( $B$ ) write( $B$ )	read( $A$ ) write( $A$ )  read( $B$ ) write( $B$ )



# Schedule 1

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
	WRITE(A)
READ(B)	
WRITE(B)	
	READ(B)
	WRITE(B)



# Schedule 1

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
	WRITE(A)
READ(B)	
WRITE(B)	
	READ(B)
	WRITE(B)



# Schedule 1

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
	WRITE(A)
READ(B)	
WRITE(B)	
	READ(B)
	WRITE(B)



# Schedule 1

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
	WRITE(A)
READ(B)	
WRITE(B)	
	READ(B)
	WRITE(B)



# Schedule 1:

Blue – NonConflicting,  
Red - Conflicting

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
	WRITE(A)
READ(B)	
WRITE(B)	
	READ(B)
	WRITE(B)



# Schedule 1:

## Black - NonConflicting( Not Swapped)

## Red - Conflicting

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
	WRITE(A)
READ(B)	
WRITE(B)	
	READ(B)
	WRITE(B)



## Schedule 2: Swapping of T1(READ(B)) and T2(WRITE(A))

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
READ(B)	
	WRITE(A)
WRITE(B)	
	READ(B)
	WRITE(B)



## Schedule 2: Swapping of T1(READ(B)) and T2(WRITE(A))

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
READ(B)	
	WRITE(A)
WRITE(B)	
	READ(B)
	WRITE(B)

Swapping of T1(READ(B)) and  
T2(READ(A))

# OR

## Schedule 3

Swapping of T1(WRITE(B)) and  
T2(WRITE(A))

T(1)	T(2)
READ(A)	
WRITE(A)	
READ(B)	READ(A)
	WRITE(A)
WRITE(B)	
	READ(B)
	WRITE(B)

T(1)	T(2)
READ(A)	
WRITE(A)	
	READ(A)
READ(B)	
	WRITE(B)
WRITE(B)	
	WRITE(A)
	READ(B)
	WRITE(B)

Any path will lead to same  
situation in next diagram



Swapping of T1(WRITE(B)) and  
T2(WRITE(A))

## OR **Schedule 4**

Swapping of T1(READ(B)) and  
T2(READ(A))

T(1)	T(2)
READ(A)	
WRITE(A)	
READ(B)	
	READ(A)
WRITE(B)	
	WRITE(A)
	READ(B)
	WRITE(B)



# Schedule 5 :

Swapping of T1(WRITE(B)) and T2(READ(A))

T(1)	T(2)
READ(A)	
WRITE(A)	
READ(B)	
WRITE(B)	
	READ(A)
	WRITE(A)
	READ(B)
	WRITE(B)



# View Serializability

Let  $S$  and  $S'$  be two schedules with the same set of transactions.  
 $S$  and  $S'$  are **view equivalent**  
if the following three conditions are met:

1. For each data item  $Q$ , if transaction  $T_i$  reads the initial value of  $Q$  in schedule  $S$ , then transaction  $T_i$  must, in schedule  $S'$ , also read the initial value of  $Q$ .
2. For each data item  $Q$  if transaction  $T_i$  executes **read( $Q$ )** in schedule  $S$ , and that value was produced by transaction  $T_j$  (if any), then transaction  $T_i$  must in schedule  $S'$  also read the value of  $Q$  that was produced by transaction  $T_j$  .
3. For each data item  $Q$ , the transaction (if any) that performs the final **write( $Q$ )** operation in schedule  $S$  must perform the final **write( $Q$ )** operation in schedule  $S'$ .

As can be seen, view equivalence is also based purely on **reads** and **writes** alone.



# View Serializability (Cont.)

- A schedule  $S$  is **view serializable** if it is view equivalent to a serial schedule.
- **Every conflict serializable schedule is also view serializable.**
- The following schedule is view-serializable but *not* conflict serializable.

$T_3$	$T_4$	$T_6$
read( $Q$ )	write( $Q$ )	
write( $Q$ )		write( $Q$ )

- Every view serializable schedule that is not conflict serializable has **blind writes**.



# Recoverability

Need to address the effect of transaction failures on concurrently running transactions.

- **Recoverable schedule** — if a transaction  $T_j$  reads a data items previously written by a transaction  $T_i$ , the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- The following schedule is not recoverable if  $T_9$  commits immediately after the read

$T_8$	$T_9$
read( $A$ )	
write( $A$ )	read( $A$ )

- If  $T_8$  should abort,  $T_9$  would have read (and possibly shown to the user) an inconsistent database state. Hence database must ensure that schedules are recoverable.



# Recoverability (Cont.)

- **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read( $A$ )		
read( $B$ )		
write( $A$ )	read( $A$ )	
	write( $A$ )	
		read( $A$ )

- If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.
- Can lead to the undoing of a significant amount of work



## Recoverability (Cont.)

- **Cascadeless schedules** — cascading rollbacks cannot occur; for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$ .
- Every cascadeless schedule is also recoverable



# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph** — a directed graph where the vertices are the transactions (names).
- We draw an arc from  $T_i$  to  $T_j$  if the two transaction conflict, and  $T_i$  accessed the data item on which the conflict arose earlier.
- We may label the arc by the item that was accessed
- A schedule is conflict serializable if and only if its precedence graph is acyclic.

