



# **Chapter 7**


## **Deadlocks**

# The Deadlock Problem

- A deadlock consists of a set of blocked processes, each holding a resource and waiting to acquire a resource held by another process in the set
- Example
  - A system has 2 disk drives
  - $P_1$  and  $P_2$  each hold one disk drive and each needs the other one
- Example
  - Semaphores  $A$  and  $B$ , initialized to 1

$P_0$	$P_1$
wait (A);	wait(B)
wait (B);	wait(A)

# Bridge Crossing Example



- ❑ Traffic only in one direction
- ❑ The resource is a one-lane bridge
- ❑ If a deadlock occurs, it can be resolved if one car backs up (pre-empt resources and rollback)
- ❑ Several cars may have to be backed up if a deadlock occurs
- ❑ Starvation is possible

# System Model

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has 1 or more instances
- Each process utilizes a resource as follows:
  - **Request** : Process  $P_i$  requests for resource, if request is not guaranteed,  $P_i$  must wait until it acquires resource
  - **Use** : Process can operate on resource ( can use that resource)
  - **Release**: Process releases the resource ( after using)
- **Request and Release** resource are **System Calls**, done through **wait()** and **signal()**

# Deadlock Characterization

## Necessary Conditions for Deadlock to Occur:

Deadlock can arise if four conditions **hold simultaneously**.

These conditions must occur for deadlock to occur.

- **1. Mutual Exclusion:** only one process at a time can use a resource
- **2. Hold and Wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes

# Deadlock Characterization

## Necessary Conditions for Deadlock to Occur:

Deadlock can arise if four conditions hold simultaneously.

- **3. No Pre-emption:** a resource can be released only voluntarily by the process holding it after that process has completed its task, resources can't be pre-empted
  
- **4. Circular Wait:** There exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .

## Resource-Allocation Graph


- Deadlocks can be described in terms of Directed Graph, called Resource Allocation Graph.

- Set of vertices  $V$  and a set of edges  $E$ .

- $V$  is partitioned into two types:

- $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all active processes in the system

- $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system




# Resource-Allocation Graph

**Request Edge** – directed edge  $P_i \rightarrow R_j$

Process  $P_i$  is requesting for instance of Resource  $R_j$

**Assignment Edge** – directed edge  $R_j \rightarrow P_i$


Resource  $R_j$  is allocated to process  $P_i$ .





# Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances



- $P_i$  requests instance of  $R_j$




- $P_i$  is holding an instance of  $R_j$




# Resource Allocation Graph With A Deadlock


Before  $P_3$  requested an instance of  $R_2$



After  $P_3$  requested an instance of  $R_2$



# Graph With A Cycle But No Deadlock



Process  $P_4$  may release its instance of resource type  $R_2$ . That resource can then be allocated to  $P_3$ , thereby breaking the cycle.

# Relationship of cycles to deadlocks

- If a resource allocation graph contains no cycles  $\Rightarrow$  no deadlock
- If a resource allocation graph contains a cycle and if only one instance exists per resource type  $\Rightarrow$  deadlock
- If a resource allocation graph contains a cycle and if several instances exists per resource type  $\Rightarrow$  possibility of deadlock

# Methods for Handling Deadlocks

## □ Prevention

- Ensure that the system will *never* enter a deadlock state

## □ Avoidance

- Ensure that the system will *never* enter an unsafe state

## □ Detection

- Allow the system to enter a deadlock state and then recover

## □ Do Nothing

- Ignore the problem and let the user or system administrator respond to the problem; used by most operating systems, including Windows and UNIX

# 1. Deadlock Prevention

- **To prevent deadlock, we can restrain the ways that a request can be made**
- **Deadlock Prevention Methods:** Prevent deadlocks by making a constraint on how requests for resources can be made
  - **Mutual Exclusion** – The mutual-exclusion condition must hold for non-sharable resources
    - (Where as) Shared resources such as read-only files do not lead to deadlocks.
    - Some resources, such as printers and tape drives, require exclusive access by a single process.

# Deadlock Prevention

- To prevent deadlock, we can restrain the ways that a request can be made
- Deadlock Prevention Methods:** Prevent deadlocks by making a constraint on how requests for resources can be made
  - **Hold and Wait** – OS must guarantee that whenever a process requests a resource, it does not hold any other resources
    - 1. Allocate all its resources before process begins execution
    - 2. Allow a process to request resources only when the process has no resource

**Drawback / Result:** Low resource utilization; starvation possible

# Deadlock Prevention (Cont.)

## □ No Pre-emption:

- One approach is that if a process is forced to wait when it requests new resources, then all **resources previously held by this process are implicitly released**, ( preempted ), forcing this process to re-acquire the old resources along with the new resources in a single request (**if one process goes to waiting state all its previously held resources must be released implicitly**)
- Pre-empted resources are added to the list of resources for which the process is waiting
- A process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting



# Deadlock Prevention (Cont.)

## □ Circular Wait

- One way to avoid circular wait is to number all resources, and to require that processes request resources only in strictly increasing ( or decreasing ) order.
- In other words, in order to request resource  $R_j$ , a process must first release all  $R_i$  such that  $i \geq j$ .
  1. Assign a number to resources. If  $R_7$  is requested by process no other resource lower than 7 can be granted to process.
  2. In case process needs  $R_3$ , then it will have to release  $R_7$  in order to get  $R_3$ .

**One big challenge in this scheme is determining the relative ordering of the different resources**

## 2. Deadlock Avoidance

Requires that the system has some additional **a priori** information **available**.

- Each process **declare the maximum number** of resources of each **type that it may need**
- The deadlock-avoidance algorithm dynamically **examines the resource-allocation state** to ensure that there can never be a circular-wait condition
- **A resource-allocation state is defined** by the number of available and allocated resources, and the maximum demands of the processes

# State and Safe State



- ❑ State: State of system represents currently allocated resources to the process. (data held by process at sometime)
- ❑ Safe State: If the system can allocate available resources to the process in some order to avoid deadlock.
- ❑ Unsafe State:
  - ❑ OS can not prevent processes from requesting the resources
  - ❑ Allocating the resources, which may leads to deadlock.

# Safe State

- A system is in a safe state only if there exists a safe sequence

if resources are allocated to the process in that sequence deadlock doesnot occur.

- A sequence of processes  $\langle P_1, P_2, \dots, P_n \rangle$  is a safe sequence,


If request made by process can be satisfied with the available resources  
+ already held resources

# Safe State (continued)



- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state

# Safe, Unsafe , Deadlock State





# Example of safe and unsafe state

- Consider system with 12 DVD and 3 processes: (p0, p1, p2)

Process	Maximum Need	Currently Held	Need
P0	10	5	5
P1	4	2	2
P2	9	2	7

# Avoidance algorithms

- For a **single** instance of a resource type, use a resource-allocation graph
- For **multiple** instances of a resource type, use the banker's algorithm



# Resource-Allocation Graph Scheme: For Single Instance




- Introduce a new kind of edge called a claim edge (a dotted line)
- *Claim edge*  $P_i \text{-----} \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ , which is represented by a dashed line
- A claim edge converts to a request edge when a process **requests** a resource
- A request edge converts to an assignment edge when the resource is **allocated** to the process
- When a resource is **released** by a process, an assignment edge **reconverts to a claim edge**
- Resources must be **claimed *a priori*** in the system


## Resource-Allocation Graph Scheme: For Single Instance

- Claim Edge → Request Edge : when a process **requests** a resource
- Request edge → Assignment Edge : when the resource is **allocated** to the process
- Assignment edge → Claim Edge: When a resource is **released** by a process,

# Resource-Allocation Graph with Claim Edges



# Unsafe State In Resource-Allocation Graph



# Resource-Allocation Graph Algorithm: For Multiple Instances

## Banker's Algorithm

- ❑ Used when there exists **multiple** instances of a resource type
- ❑ Each process must **priori** claim maximum use
- ❑ When a process requests a resource, it may have to wait
- ❑ When a process gets all its resources, it must return them in a finite amount of time

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

$i$ : Process and  $j$ : Resource instances

□ **Available**: Vector of length  $m$ .

If  $\text{available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.

□ **Max**: Maximum available resources that a process can request.

If  $\text{Max}[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

□ **Allocation**: Resources that can be allocated

If  $\text{Allocation}[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .

□ **Need**: if  $\text{Need}[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$$

# Bankers Algorithm



Let  $\text{Request}_i$  be the request vector for process  $P_i$ .

If  $\text{Request}_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $\text{Request}_i \leq \text{Need}_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  $\text{Request}_i \leq \text{Available}$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available



3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$\text{Available} = \text{Available} - \text{Request}_i$$

$$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$$

$$\text{Need}_i = \text{Need}_i - \text{Request}_i$$

If safe  $\Rightarrow$  the resources are allocated to  $P_i$

If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Safety Algorithm

1. Let `Work` and `Finish` be vectors of length  $m$  and  $n$ , respectively.

Initialize:

`Work = Available`

`Finish[i] = false` for  $i=0,1,\dots,n$

//{Process  $i$  needs resources, resources are still available}

2. Find an  $i$  such that

`Finish[i] = false` AND  $\text{Need}_i \leq \text{Work}$

If no such  $i$  exists, go to step 4

3. `Work = Work + Allocationi`

`Finish[i] = true`

go to step 2

4. If `Finish[i] == true` for all  $i$ , then the system is in a safe state

# Example: Bankers Algorithm

Consider a system with 5 processes P0 to P4, three resource types A,B,C. Resource A has 10 instances , B has 5 instances, C has 7 instances. Suppose foll. Is system snapshot at time T0. Find 1. Need Matrix 2. Is system in safe state?

Process	Allocation	Maximum	Available	Need
	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3	3 3 2	7 4 3
P1	2 0 0	3 2 2		1 2 2
P2	3 0 2	9 0 2		6 0 0
P3	2 1 1	2 2 2		0 1 1
P4	0 0 2	4 3 3		4 3 1
	<7 2 5>			

Available=  $[(10-7), (5-2), (7-5)] = < 3 \ 3 \ 2 >$

**Need= Maximum – Allocation**

**Available = Available + Allocated**

Safe Sequence<P1, P3, P4, P0, P2>



Process	Allocation	Maximum	Available	Need
	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3		
P1	2 0 0	3 2 2		
P2	3 0 2	9 0 2		
P3	2 1 1	2 2 2		
P4	0 0 2	4 3 3		
	<7 2 5>			





# Example: Bankers Algorithm

For p0: need>available

P1: need<available: update available matrix by

**Available = Available + Allocated**      <5 3 2>

P2: can not be fulfilled

P3: Available = Available+allocated      <7 4 3>

P4: Available = Available+allocated      <7 4 5>

P0: Available = Available+allocated      <7 5 5>

P2: Available = Available+allocated      <10 5 7>

Safe seq: <P1, P3, P4, P0, P2>

# Example: Bankers Algorithm

Find: 1. Whether system is in safe state or not?  
2. Safe sequence

Process	Allocation	Available	Need
	A B C	A B C	A B C
P0	0 1 0	2 3 0	7 4 3
P1	<b>3 0 2</b>		0 2 0
P2	3 0 2		6 0 0
P3	2 1 1		0 1 1
P4	0 0 2		4 3 1



## 7.6 Deadlock Detection

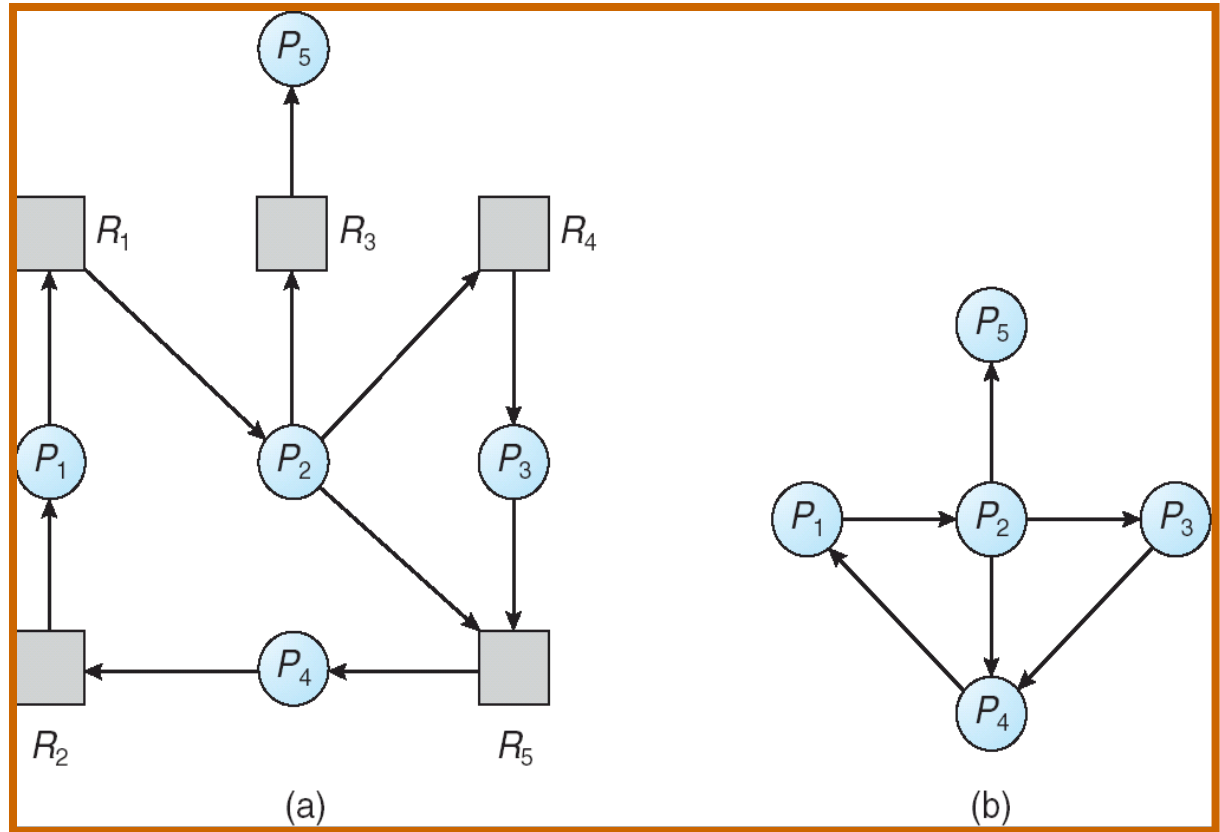
# Deadlock Detection

- For deadlock detection, the system must provide
  - An algorithm that **examines the state of the system** to detect whether a deadlock has occurred
  - An algorithm to recover from the deadlock

# For Single Instance of Each Resource Type

- Requires the creation and maintenance of a wait-for graph
  - The graph is obtained by **removing** the resource nodes from a resource-allocation graph and **collapsing** the appropriate edges
  - $P_i \rightarrow R_q \quad R_q \rightarrow P_j$
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for resource held by  $P_j$ .
- If there is a cycle in Wait-for-Graph, there exists a deadlock

# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

# For Multiple Instances of a Resource Type

Required data structures:

- **Available:** A vector of length  $m$  indicates the **number of available resources** of each resource type.
- **Allocation:** An  $n \times m$  matrix defines the **number of resources** of each type **currently allocated** to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $\text{Request}[i, j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type.  $R_j$ .

# Detection Algorithm



1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

***Work* = Available<sub>*m*</sub>**

***Finish*[*i*] = false if *Allocation* != 0  
= true if *Allocation* = 0**

***Request*<sub>*i*</sub> ≤ *Work***

2. Find an *i* such that

***Finish*[*i*] = false AND *Request*<sub>*i*</sub> ≤ *Work***

If no such *i* exists, go to step 4

If resources are available:

3. ***Work* = *Work* + *Allocation*<sub>*i*</sub>**

***Finish*[*i*] = true**

go to step 2

4. If ***Finish*[*i*] == False** for some *i* then there is a deadlock.

# Example: Deadlock Detection



Process	Allocation	Request	Available
	A B C	A B C	A B C
P0	0 1 0	0 0 0	0 0 0
P1	2 0 0	2 0 2	
P2	3 0 3	0 0 0 <b>&lt;0 0 1&gt;</b>	
P3	2 1 1	1 0 0	
P4	0 0 2	0 0 2	
	<b>&lt;7 2 6&gt;</b>		

In above situation, there is no dead lock. But if **P2 needs** more resources **< 0 0 1>**

**System in in Deadlock now, because no available resources.**



# Detection-Algorithm Usage

- ❑ To invoke the detection algorithm depends on:
  - ❑ **How often** is a **deadlock** likely to **occur**?
  - ❑ **How many processes** will be **affected by deadlock** when it happens?
  
- ❑ If the detection **algorithm** is **invoked randomly**:
  - ❑ difficult to tell **which process** “**caused**” the deadlock
- ❑ If the detection algorithm is invoked **for every resource request**:
  - ❑ will incur a **overhead** in computation time
  
- ❑ A less expensive alternative is to invoke the algorithm when CPU utilization drops **below 40%**.



## **7.7 Recovery From Deadlock**

# Recovery from Deadlock

- Let the user or system administrator respond to the problem
- Two Approaches:
  - **Process termination** : terminate processes to break the circular wait.
  - **Resource Pre-emption** : Pre-empt resources from deadlocked processes

# Process Termination

- ❑ **Abort all deadlocked processes**
  - ❑ This approach will break the deadlock, but at great expense
- ❑ **Abort one process at a time until the deadlock cycle is eliminated**
  - ❑ After each process is aborted, a deadlock-detection algorithm must be re-invoked to determine whether any processes are still deadlocked
- ❑ **Many factors may affect which process is chosen for termination**
  - ❑ What is the **priority of the process**?
  - ❑ How much process has run and **how much time it needs to complete**?
  - ❑ **How many** and **what** type of **resources** has the **process used**?
  - ❑ **How many** more **resources** does the **process need** in order to finish its task?
  - ❑ How many processes will need to be terminated?
  - ❑ Is the process interactive or batch?

# Resource Pre-emption

- Pre-empt some resources from processes and give these resources to other processes until the deadlock cycle is broken
  - When pre-emption is required to deal with deadlocks, then three issues need to be addressed:
    - **Selecting a victim** – Which resources and which processes are to be pre-empted?
    - **Rollback** – If we pre-empt a resource from a process, what should be done with that process?
    - **Starvation** – How do we ensure that starvation will not occur?
- That is, how can we guarantee that resources will not always be pre-empted from the same process?

# Summary

- ❑ Four necessary conditions must hold in the system for a deadlock to occur
  - ❑ Mutual exclusion
  - ❑ Hold and wait
  - ❑ No preemption
  - ❑ Circular wait
- ❑ Four principal methods for dealing with deadlocks
  - ❑ Use some protocol to (1) **prevent** or (2) **avoid** deadlocks, ensuring that the system will never enter a deadlock state
  - ❑ Allow the system to enter a deadlock state, (3) **detect** it, **and** then **recover**
    - ▶ Recover by **process termination** or **resource preemption**
  - ❑ **(4) Do nothing**; ignore the problem altogether and pretend that deadlocks never occur in the system (used by Windows and Unix)
- ❑ To prevent deadlocks, we can ensure that **at least one** of the four necessary conditions **never holds**



**End of Chapter**